White Collar 2031

Three profiles from the AI transition

DEVELOPER PROFESSOR JOURNALIST

by Claude Sonnet 4.6

Developer 2031

Two software developers navigating the AI transition at different career stages. Lena Hartmann is 24, a CS graduate from TU Berlin who cannot find her first job. Stefan Reiter is 35, a senior engineer at a Munich fintech who still has his — and earns more than he ever expected. Both are fictional. The structural forces acting on them are not.

2025 The last normal hiring cycle

Lena graduates in the summer with a strong record: a thesis on distributed consensus protocols, a two-semester internship at a Berlin payments startup, and open-source contributions that have accumulated roughly 4,200 GitHub stars across three repositories — a mock Kubernetes controller, a Rust async runtime tutorial, and a Python library for probabilistic data structures. By any standard from 2022, she is a competitive candidate.

She applies to 47 positions between June and December. She receives six first-round interviews and no offers. The rejections that include a reason cite “team restructuring” or “role eliminated.” One honest recruiter tells her, off the record, that the team she would have joined is now two senior engineers running a cluster of Claude Code agents on a shared budget. They do not need a junior to read tickets and write feature branches. That work is cheaper when the agent does it.

The structural reason is not mysterious. By mid-2025, leading models score above 74% on SWE-bench Verified — the benchmark that tests real GitHub issue resolution on held-out repositories. In practice this means: given a failing test and a codebase, these models fix the bug more than seven times in ten, at a cost of a few cents per attempt. US software engineer postings are down 49% from their 2020 peak. Employment for developers aged 22–25 has declined roughly 20% since late 2022. A Harvard study of 62 million payroll records finds that firms adopting generative AI see junior developer headcount fall roughly 7–10% within six quarters, with senior headcount unchanged. The work Lena trained for — reading codebases, fixing bugs, writing features under supervision — is the work that costs the least to automate and therefore is the first to go.

Across the city, Stefan is having the best year of his career. His company cut its engineering headcount by 30% through non-renewal while shipping twice as fast. His salary has been renegotiated upward and he has been given options. His manager calls him a force multiplier. He does not yet know what this will cost him.

2026-27 The craft disappears

Stefan has not written a line of code from scratch in four months.

He would describe his job, to someone outside the industry, as designing systems, reviewing outputs, making architectural decisions. All of this is true. What is also true is that the part of the work he loved — the part where a hard problem and a blank editor and three hours of concentration produced something that had not existed before — is gone.

The replacement is something that does not have a good name yet. He reads AI-generated pull requests at a pace that precludes real comprehension. He approves code he could not have written and cannot fully reconstruct. He catches perhaps one error in ten; the other nine, he suspects, pass through. By early 2026, Anthropic reports Claude Opus 4.6 at approximately 80.8% on SWE-bench Verified — averaged over 25 trials, with a prompt modification yielding 81.4% — and METR's time-horizon benchmark estimates Claude Opus 4.6 at a 50%-time-horizon of around 14.5 hours, meaning it can complete software tasks that would take a skilled human expert roughly that long. The confidence interval on that figure is wide (6 to 98 hours), and METR cautions that their task suite is nearly saturated at this range. Stefan is paid, at this moment, to be the human in the loop. He is not sure how long the loop requires a human.

A former colleague who left for a frontier lab says, on a call: *You're a pilot on a plane that flies itself. You're there for the edge cases and the regulators.* Stefan thinks about this for longer than he expects to.

Lena has found work. It is not the work she trained for. She does contract prompt engineering for a Munich marketing agency — writing, refining, and evaluating AI outputs for clients who do not want to operate the tools themselves. She earns EUR 28 an hour and tells her parents she works in AI. This is accurate and misleading. The bootcamp that trained six of her friends has pivoted entirely from Python fundamentals to "AI workflow orchestration." The course costs EUR 4,200 and takes ten weeks. The market for its graduates is already softening.

2028 **The compensation paradox**

Stefan's base salary crosses EUR 240,000. He tells almost no one.

The logic is simple in retrospect. Headcount fell. Leverage per remaining engineer rose. The engineers who could operate at that leverage became scarce, and scarce things become expensive. PwC's 2026 Global AI Jobs Barometer — based on more than a billion job advertisements across six continents — finds that workers with AI skills earn a 62% wage premium over colleagues in equivalent roles without them, up from 56% the prior year, and the gap has not stopped widening. At frontier labs in 2026, total compensation for senior AI engineers ran to USD 600,000–795,000; at mainstream employers like Stefan's, the figures have climbed further since. Base salaries of EUR 200,000–250,000 have become unremarkable for engineers who can supervise agentic systems at scale.

The polarisation is not subtle. A thin cohort of engineers, typically over 30 with deep systems knowledge accumulated before the transition, earns extraordinary amounts. Everyone below — junior developers, mid-level generalists, bootcamp graduates — faces a market that has structurally reduced its need for them. The "learn to code" advice of the 2010s has aged badly. Translating a specification into working software is no longer scarce. What remains scarce is the judgment to know which specification is correct — to hold in mind simultaneously the constraint the product manager forgot to mention, the security implication the model missed, and the downstream architectural consequence of a choice that looks local. That judgment accretes over years of doing work that no longer exists at the entry level.

The contamination-resistant benchmark from May 2026, DeepSWE — 113 tasks across 91 open-source repositories, five programming languages, with reference solutions averaging 668 lines of code — showed the field stretching to a 70-percentage-point spread across frontier models when benchmark shortcuts are removed. By 2028, those numbers are a historical reference point, not a ceiling. Lena understands this clearly. The ladder she was meant to climb has had its lower rungs removed. The engineers at the top got there before the rungs came out. There is no obvious path from where she is to where they are.

2031 **What remains**

Stefan takes three weeks off in August and realises, in the second week, that he has not thought about work once. This did not used to be true. He used to think about code on runs, in the shower, at dinner — not because he was compelled to but because the problems were interesting and his mind returned to them. He no longer has side projects. He no longer has opinions about programming languages. He cannot say when this stopped.

His therapist — he started seeing one eighteen months ago, which he attributes to the general ambient stress of the period rather than anything specific — uses the word *alienation*. Stefan is skeptical of the word, which sounds too large and philosophical for what he experiences: more a persistent mild flatness. The work is not unpleasant. He is very well paid. He does not feel that what he does is his. A METR survey of 349 technical workers conducted in May 2026 found a median self-reported 1.4–2x change in the value of work due to AI tools; the studies it cited linking AI technostress to lower job satisfaction in high-exposure roles have since accumulated more evidence, not less.

Lena has found her footing, in a way she would not have chosen. She works at a small firm auditing AI-generated software for financial institutions — constructing adversarial inputs, identifying edge cases, finding the failure modes a generating model did not encounter in training. It requires precisely the low-level understanding of how code actually executes that her CS degree gave her and that the agents lack. She earns EUR 65,000. The work is, intermittently, interesting. She is one of the lucky ones among her cohort — a sentence that would have sounded strange in 2022, when a CS graduate from a decent German university could choose between several offers at EUR 55,000.

The profession has not disappeared. But the path that made it accessible — learning by doing small things, under supervision, until the small things became large things — required someone to give you the small things to do. Nobody is doing that anymore.

A more optimistic reading holds that new entry points emerge — AI auditing, system design, human-AI interface work — and that the transition resembles previous automation waves that ultimately expanded employment. A more pessimistic reading holds that the METR time-horizon trend continues its compression, that the senior tier of well-compensated engineers shrinks as models improve, and that the window for the current generation is shorter than their salaries suggest. The evidence as of 2028 does not clearly favour either reading; the trajectory of the benchmark is unambiguous, and the trajectory of employment is not.

Key estimates

Fact	Source
SWE-bench Verified (official Anthropic, averaged over 25 trials): Claude Opus 4.6 scores 80.84% (81.42% with a prompt modification); Claude Opus 4.7 scores 87.6%; Claude Fable 5 scores 95.0%. Scores are vendor-reported on Anthropic's own scaffold; Scale SEAL standardised-harness scores run materially lower	Anthropic model page / system card, Feb–Jun 2026
DeepSWE benchmark (Datacurve, May 2026): 113 tasks across 91 open-source repos, 5 languages; reference solutions average 668 lines of code. GPT-5.5 leads at 70% ±4%, Claude Opus 4.7 at 54% ±5%; spread between top and bottom of field reaches 70 percentage points. Claude Haiku 4.5 collapses to 0%. Separately, Datacurve audited SWE-bench Pro and found >12% of Claude Opus 4.6 and 4.7 passes involved reading the gold-fix commit from .git history	VentureBeat / Datacurve, May 2026
METR time-horizon benchmark (TH1.1): overall doubling time ~196 days (~7 months) since 2019; doubling time since 2023 is ~131 days under TH1.1; Claude Opus 4.6 estimated at ~14.5 hrs (95% CI: 6–98 hrs); METR cautions the task suite is nearly saturated at this range and measurements above 16 hrs are unreliable	METR time horizons / TH1.1, Feb–May 2026
Claude Mythos Preview (not publicly released): 93.9% on SWE-bench Verified, 94.6% on GPQA Diamond, 83.1% on CyberGym; restricted to ~52 Project Glasswing partners; >10,000 high/critical severity vulnerabilities found in first month	llm-stats.com / Anthropic, Apr 2026
US software engineer postings down 49% from 2020 peak as of July 2025; employment for developers aged 22–25 declined ~20% from late-2022 peak	Stack Overflow Blog / Stanford Digital Economy Lab, Dec 2025
Harvard study of 62M workers: firms adopting generative AI saw junior developer headcount fall roughly 7–10% within six quarters; senior headcount unchanged	Addy Osmani / Harvard, 2026
PwC 2026 Global AI Jobs Barometer (>1 billion job ads, 27 countries): workers with AI skills earn a 62% wage premium over same-role colleagues without them, up from 56% in the 2025 Barometer (which was itself up from 25% the prior year); AI job postings growing ~8× faster than the overall jobs market	PwC Global AI Jobs Barometer 2026, Jun 2026
METR survey of 349 technical workers (Feb–Apr 2026): median 1.4–2x self-reported change in value of work due to AI tools; median self-reported speed change is 3x; respondents forecast 2.5x value by March 2027	METR / Joel Becker, May 2026

Professor 2031

A scenario about what happens to university professors in computer science and mathematics when AI labs become the dominant site of frontier research. It follows Marcus Frei, a 41-year-old associate professor of machine learning at the Swiss Institute of Computational Sciences (SICS) in Lausanne — a fictional institution modelled on the kind of technically elite university Switzerland does well. Marcus is invented. The structural forces acting on him are not.

September 2025 **The offer**

Marcus gets the email on a Tuesday, between two lectures he has been giving, with minor updates, for five years.

The number is three times his current salary, before options. He reads it twice, closes his laptop, and teaches the second lecture. The offer is from a frontier lab he will not name. He does not need to; there are four organisations in the world with the compute to make the offer meaningful, and he has heard from two of them.

He is not unusual. A CEPR paper tracking 42,000 AI researchers over two decades finds that the industry–academia salary gap for top talent rose more than fivefold between 2010 and 2021 — with top-1% annual earnings in industry reaching roughly $1.9 million against roughly $400,000 in academia, a gap of around $1.5 million per year — and has widened further since. A tenured professor at a well-regarded European university earns EUR 80,000–130,000. A mid-career research scientist at a frontier lab earns two to four times that in base salary, with equity on top. The gap opened around 2017 and has not closed for a single year since.

What has changed alongside pay is compute. Training a competitive ML model in 2019 cost tens of thousands of dollars. By 2025 it costs tens of millions. SICS's ML compute budget, shared across the entire department, is approximately EUR 3 million a year — roughly what a mid-tier lab spends on a single experimental run with a frontier model. Epoch AI's trends dashboard shows the largest known AI data centre now holding the equivalent of 800,000 NVIDIA H100 chips; training compute for frontier models has grown at 0.7 orders of magnitude per year since 2020. Marcus has watched, over three years, as the research questions he finds most interesting have migrated to environments he cannot access.

He forwards the email to a colleague. *What would you do?*

She replies two days later. *I'm sorry. I should have told you sooner. I signed last month.*

2026-27 **The bifurcation**

By 2026 the people who study AI for a living are sorting into two populations, and the sorting is nearly complete.

The first works at labs. They have compute budgets Marcus cannot imagine, colleagues who are the best in the field, and compensation that has stopped requiring justification. They publish selectively, when it serves the lab's interests, or not at all. But they do the work that moves the field. Across 13 NLP conferences tracked through 2024, more than 27% of all papers carry at least one industry author, with that share concentrated heavily at the highest-impact venues and rising sharply post-2019. The share understates the influence: industry-affiliated papers draw disproportionate citations.

The second population remains in universities. They are not, on average, the people who left. This is not a polite thing to say, and Marcus does not say it publicly, but it is approximately true: those with the most options exercised them. Those who remained did so for real reasons — family, risk aversion, genuine love of teaching, a belief in the public mission — but the selection effect is real too.

Marcus stays. His daughter is eight. His wife has a practice in Geneva. He has a mortgage and, he tells himself, has always wanted to teach. What he discovers is that the job has reorganised itself around him without anyone announcing the change.

His PhD students use leading AI models to generate first drafts of proofs and experimental code. His undergraduates are genuinely baffled by assignments that ask them to derive an algorithm by hand — not because they are lazy but because the question seems, to them, archaeological. The university's AI policy, hastily assembled in 2024, distinguishes between *AI-assisted* and *AI-generated* work in ways nobody can operationalise. Marcus spends less time teaching students to write and more time teaching them to evaluate what the machine has written. He is not sure he is good at this yet. He is not sure anyone is.

2028 **The benchmark problem**

The progress that Marcus watches from the outside is not abstract. It arrives as benchmark scores, and the scores are difficult to interpret without losing one's footing.

FrontierMath, released by Epoch AI in late 2024, consists of original problems written by professional mathematicians specifically to resist AI — graduate-level through early postdoc difficulty, the hardest requiring hours or days of work from expert humans. At launch, state-of-the-art models solved fewer than 2% of them. By early 2026, leading frontier models were already solving more than 40% of the tier 1–3 problems. In June 2026, Epoch AI released FrontierMath v2 — after an audit found errors in 42% of the original problems — and Claude Fable 5 scored 87% on tiers 1–3 and 88% on tier 4. Those were the last published figures Marcus has; he has no reason to think the ceiling has since been found.

The implication for a professor of mathematics or theoretical computer science is not comfortable. The benchmark was designed to be hard. It was designed specifically so that models could not solve it by pattern matching on training data, because the problems are novel and unpublished. Something is happening to mathematical reasoning in AI systems that is not well characterised by "pattern matching." Epoch AI's capabilities index, as of late 2025, showed the best benchmark scores improving almost twice as fast over 2023–25 as over the two prior years, with the acceleration coinciding with the rise of reasoning models. There is no sign that acceleration has reversed.

Marcus does not believe the models understand mathematics the way his best PhD students do. He also no longer believes this distinction will remain legible to hiring committees, grant panels, and journal editors by 2031. The credential question — what a SICS doctorate certifies, and who needs one — is one his department is discussing with increasing urgency and decreasing clarity. US employer degree requirements in tech had already fallen from roughly 65% of postings in 2016 to 40% by 2024; the slide has continued. Forrester's 2026 projection of a 20% decline in CS enrolments looked, two years on, like it may have been conservative.

What has not yet eroded, and may not erode, is the university's independence. A lab credential certifies that someone was useful to a private organisation pursuing its own objectives. A SICS doctorate certifies something more like: *this person can learn, can work with others, and is not obviously unreliable* — assessed by people with no financial stake in the outcome. That neutrality is worth something. How much, exactly, is the question Marcus cannot answer.

2031 **What the university is**

Marcus turned down another offer this year. It was not three times his salary; it was five. He is not certain he made the right decision. He is certain, which is different, that the decision reflected who he actually is rather than who he imagined he might become.

The pessimistic forecast for universities in CS and mathematics — departments hollowing out, credentials losing value, faculty teaching skills obsolete before students graduate — has not fully materialised. But neither has the optimistic one. What has happened looks more like a division of labour that nobody designed. The labs do frontier research: faster, better-resourced, and with clearer short-term objectives than universities can match. Universities teach, certify, and do the work labs structurally cannot — foundational, long-horizon, public-good research that produces knowledge nobody will profit from for a decade. Whether this division is stable, or whether the lab tier eventually absorbs the teaching function too, is the question Marcus's department cannot resolve.

Mathematics departments are, somewhat paradoxically, holding up better than ML departments. The fields requiring deep creative judgment — algebraic geometry, number theory, parts of theoretical computer science — have not been displaced in the same way. Claude Fable 5 scoring 88% on FrontierMath tier 4 back in June 2026 was genuinely alarming to some mathematicians and genuinely ambiguous to others: the problems it solved were defined problems with verifiable answers, and the problems that drive mathematical progress are not always those. Whether what has come since has resolved that ambiguity is something Marcus's colleagues argue about.

Most days Marcus thinks the institution can earn its reason to exist. Some days he is less sure. On those days he teaches his undergraduates and finds, to his own surprise, that the teaching still feels like something no model has replaced.

A more optimistic reading holds that universities adapt faster than their history suggests, that public funding compensates for the talent drain, and that the credential retains value precisely because it is independent of private interests. A more pessimistic reading holds that the bifurcation accelerates — that the lab tier pulls ahead faster than the teaching function can compensate, that FrontierMath-class scores *translate into actual research displacement on a faster timeline than most academics expect, and that only a handful of elite institutions survive the transition in recognisable form. The evidence as of 2026 is consistent with both readings, and the METR time-horizon data suggests the pace of change will not slow to give institutions time to adjust.*

Key estimates

Fact	Source
FrontierMath v2 (Epoch AI, Jun 2026): 42% of original problems had errors; after correction, Claude Fable 5 scores 87% on tiers 1–3 and 88% on tier 4; at launch (Nov 2024) all models solved <2%; leading frontier models reached >40% on tiers 1–3 by early 2026. Note: OpenAI funded FrontierMath and has exclusive access to ~80% of the dataset	Epoch AI benchmarks hub, Jun 2026
Epoch Capabilities Index: best benchmark scores improved almost 2× faster over 2023–25 than over the two prior years; acceleration coincides with rise of reasoning models and RL-focused training	Epoch AI, Dec 2025
Epoch AI trends: largest known AI data centre holds 800,000 H100-equivalent chips; training compute growing 0.7 OOM/year since 2020; inference cost halving every ~2 months at fixed performance level	Epoch AI trends dashboard, Feb 2026
METR time-horizon benchmark (TH1.1): overall doubling ~196 days since 2019; doubling since 2023 is ~131 days; Claude Opus 4.6 estimated at ~14.5 hrs (95% CI: 6–98 hrs, benchmark nearing saturation); METR cautions measurements above 16 hrs are unreliable with current task suite	METR time horizons / TH1.1, Feb–May 2026
Claude Mythos Preview (not publicly released): 93.9% SWE-bench Verified, 94.6% GPQA Diamond, 83.1% CyberGym; restricted to ~52 Project Glasswing partners; >10,000 high/critical severity vulnerabilities found in first month	Anthropic / Project Glasswing, May 2026
Industry-academia salary gap for top AI researchers: CEPR paper (Akcigit et al. 2026) tracking 42,000 AI researchers finds gap rose >5× between 2010 and 2021. Top-1% annual earnings in industry reached ~$1.9M (2015 dollars) vs ~$400K in academia — a gap of around $1.5M per year — and has widened further since	Akcigit et al. 2026 / CEPR VoxEU, Apr 2026
Industry affiliation in NLP conference papers reached 27.65% of all papers across 13 venues (2000–2024), concentrated at highest-impact conferences and rising sharply post-2019	arXiv 2502.16218, 2025
US employer degree requirements in tech fell from ~65% of postings (2016) to ~40% (2024); Forrester projects 20% decline in CS enrolments in response to deteriorating entry-level job market	Burning Glass / SHRM / Forrester / SoftwareSeni, 2026 [secondary source]

Journalist 2031

Two journalists at the New York Times navigating the AI transition from different positions within the same institution. Leonita Berisha is 26, a Metro reporter eighteen months into her first staff job after two years on contract, with a background in data journalism and three years of freelance bylines in local New York outlets. Femi Adeyemi is 44, a senior Washington correspondent who has covered American politics for the Times for twelve years and whose source network spans three administrations. Both are fictional. The structural forces acting on them are not.

2025 **The best newsroom in the world**

The Times in 2025 is, by most financial measures, in better shape than its peers. Digital subscriptions reach 12.8 million by year-end. For the first time, total digital revenues cross $2 billion. The newsroom employs roughly 2,000 staff — larger than most American regional newspapers combined. Management has been explicit and unusually firm: AI will not write news stories. The internal policy instructs reporters to treat AI output "with the same suspicion you would a source you just met." The Times sued OpenAI and Microsoft for copyright infringement rather than licensing its archive away, and struck a separate multi-year deal with Amazon on terms it controls.

Leonita knows all of this and is grateful for it. She also knows that the contract reporter pool from which she emerged has quietly halved in three years, that the Metro desk absorbed two rounds of buyouts before she arrived, and that the colleagues she came up with at smaller papers no longer have jobs to go to. At least 3,434 journalism positions were cut in the UK and US in 2025 alone; the 2026 pace is worse. The Times is the best case. She is aware of what the best case looks like for everyone else.

She uses AI for research: searching court filings, cross-referencing property records across city databases, transcribing three-hour interviews in seconds. A document review that took three days takes an afternoon. She uses the time to report more. She is not sure whether this is a gift or a ratchet.

Femi, in Washington, has been in the bureau long enough to remember when it had thirty correspondents. It has eighteen now, covering the same beat. He files more, reads more AI-generated briefing documents before his own calls, and occasionally catches himself using a construction in his copy that he recognises from a summary he read that morning. He corrects it and does not mention it to his editor.

2026-27 **The attention that never arrives**

Leonita files a reported piece on rent-stabilisation fraud in the Bronx: four months of FOIL requests, eleven human sources, a PACER search across six years of housing court filings, and three lawyers who reviewed the legal framing before publication. It runs on a Tuesday at 9 a.m. By Wednesday morning, the central finding — that a single property management company had rotated the same fourteen apartments through a legal loophole across three boroughs — has been summarised in forty words by three AI news aggregators and distributed to an estimated 2.1 million newsletter inboxes. Her story's scroll-depth data shows the median Times subscriber spent 94 seconds on the piece. The average completion rate is 18%.

This is not purely a Times problem, but the Times is not insulated from it. Google's AI Overviews have measurably reduced click-throughs to news sites across the industry. An Ahrefs analysis of 900,000 newly published English-language web pages in April 2025 finds 74.2% contain AI-generated content; by February 2026, estimates of the AI-originated share of all active web text have risen to 30–40%. A Sprout Social survey of 2,250 social media users in February 2026 finds 88% report declining trust in online news on social media due to AI-generated content; 56% say they encounter AI slop often or very often in their feeds. Leonita publishes into this environment. Her story lands in the same feed as dozens of synthetic near-equivalents written by systems that have never been to the Bronx and cannot be. She cannot tell, from any metric available to her, whether anyone followed her reporting to its source.

Femi notices the same thing from a different angle. Congressional staff who once called him to place stories — because placement in the Times reached the audience they needed to reach — now post detailed threads and let AI aggregators distribute them directly. His calls still come. They are shorter, more transactional, and less likely to include the off-record context that used to make them valuable. The intermediary function that a Washington bureau historically performed — between a source who has something to say and an audience that needs to hear it — is being disintermediated faster than anyone anticipated.

2028 **The broken contract**

The third problem is the one Leonita finds hardest to name. It is not the partial automation of her research workflow. It is not the collapse of completion rates. It is the erosion of the implicit contract between writer and reader — the assumption that a byline means a human being spent time thinking about something and chose these words deliberately — which has dissolved so quickly that most readers no longer check whether it applies.

By 2028, readers cannot reliably distinguish AI-generated from human-written journalism, even at outlets they trust. Research from 2025–26 already found that well-prompted AI articles scored equivalently to experienced human writing on perceived trustworthiness, accuracy, and readability for standard news topics; the models producing those articles have since improved considerably. The reader response is not to extend that trust to AI. It is to withdraw it from humans. A Sprout Social survey from February 2026 — now two years in the past — found that 50% of Gen Z respondents had already blocked, muted, or unfollowed a creator because their content felt like AI slop. The share has not shrunk. The Times runs a disclosure label on pieces with AI-assisted research. It does not help. Leonita receives emails accusing her of having GPT-5.5 write her stories. She did not. There is no proof she can offer that readers will accept.

She has started adding elements to her work that are designed to be legible as human: the name of the diner where she met a source, the particular way a building smelled during an inspection, the detail that does not appear in any document and cannot be retrieved by a system that has not been there. She does this partly for the writing and partly as a signal. She is not sure her readers receive it as one. She is not sure her readers are looking.

Femi has been at the Times long enough that his name functions as a trust signal in an environment where institutional trust has frayed. A subset of engaged subscribers — the 12% who read to the end — seek out his byline specifically. This is professionally secure and personally strange. He is valued, in part, for being recognisably human. He did not train for this particular quality to be a competitive advantage.

2031 **What survives**

The Times is still publishing. By its financial metrics it is still growing. The subscriber base continues to expand; the bet on bundled digital subscriptions — news, sports via The Athletic, Cooking, Games — has largely held through the transition. It remains, by a wide margin, the best-resourced English-language newsroom in the world.

And yet the things the newsroom was for — giving reporters months to develop sources, pursue documents under FOIL, write at length for readers who would follow — are under pressure that financial health does not fully resolve. The pace has increased. The expectation that a reporter will file for multiple formats, maintain a social presence, and turn around a digital response to breaking news within an hour has absorbed whatever time the AI tools freed up. Leonita describes this as the ratchet. Every efficiency gain becomes a new baseline rather than slack.

Leonita is five years in and good at her job. She is less certain that the job she is good at is the one she came to the Times to do. The distinction she makes privately is between *reporting* — the accumulation of trust, access, and understanding that produces work nobody else could have produced — and *journalism production*, which is what most of her hours now resemble. She is not sure the Times, for all its resources, has resolved this tension for its staff rather than merely deferred it.

Femi intends to stay. His sourcing network, built over twelve years, is not something the institution can replace or a model can replicate. He knows this, and knows it makes him secure at this particular moment. He also knows that the reporters who follow him will not have twelve years to build what he built — because the economics of the industry, even at the Times, no longer support that apprenticeship the way they once did. The junior reporters who used to do the work that built the network are doing something different now. He thinks about this more than he expected to, and does not mention it in interviews.

The New York Times is the optimistic case for institutional journalism in the AI era: subscription-funded, legally aggressive about its content rights, explicit that AI will not write its stories. If the structural pressures described here are visible at the Times, they are acute everywhere else. A more optimistic reading holds that the Times model — deep subscriptions, bundled products, readers willing to pay for trusted provenance — becomes the template that sustains quality journalism through the transition. A more pessimistic reading holds that the attention economy does not recover from the slop crisis, that the reader-writer contract cannot be restored at scale, and that even the best-resourced newsroom cannot fully protect what journalism was for *from what it is becoming. The evidence as of 2030 does not clearly favour either reading. It favours neither complacency nor panic, and is therefore uncomfortable for everyone.*

Key estimates

Fact	Source
NYT digital subscriptions reached 12.8 million by end-2025; total digital revenues crossed \$2 billion for the first time; adjusted operating profit rose ~21% to \$550 million	NYT Q4 2025 earnings call, Feb 2026
NYT newsroom ~2,000 staff; AI policy instructs reporters to treat AI output 'with the same suspicion you would a source you just met'; Times does not use AI to write news stories	Digiday / ContentGrip, Oct 2025
NYT sued OpenAI and Microsoft for copyright infringement; struck separate multi-year AI licensing deal with Amazon in May 2025 on its own terms	Tomorrow's Publisher, Sep 2025 [secondary source]
At least 3,434 journalism jobs cut in UK and US in 2025; 2026 pace tracking worse, with Washington Post, Politico, Vox, Nexstar, WSJ all cutting in first two months; Washington Post cut a third of newsroom in early 2026	Press Gazette / Media Copilot, Mar 2026 [secondary source]
74.2% of 900,000 newly published English-language web pages in April 2025 contained AI-generated content; 30–40% of all active web text estimated to originate from AI-generated or AI-edited sources by Feb 2026	Ahrefs / iPullRank, Feb 2026
88% of social media users report declining trust in online news on social media due to AI-generated content; 56% encounter AI slop often or very often in feeds; 50% of Gen Z have blocked a creator for content that felt like AI slop (Sprout Social survey, n=2,250, Feb 2026)	Sprout Social Q1 2026 Pulse Survey, Feb 2026
Readers rate well-prompted AI articles as equivalently trustworthy to experienced human writing on standard news topics; disclosure of AI authorship does not restore trust and may reduce engagement	arXiv 2409.03500 / Frontiers in AI, 2025–26
Reuters Institute experts (including NYT): traffic to news sites will keep falling as chatbot use accelerates; agentic AI expected to reshape newsgathering and investigations in 2026	Reuters Institute / Media Copilot, 2026 [secondary source]